

# Solving the 8-Puzzle Problem with A\*

## An Implementation in Python

Yule Liu

HKUST(GZ) - UFUG 1601

July 8, 2025

# What is the 8-Puzzle Problem?

The 8-puzzle is a classic sliding puzzle that consists of a 3x3 grid with 8 numbered tiles and one empty space.

**Objective:** Rearrange the tiles from a given initial configuration to a target goal configuration by sliding tiles into the empty space.

Initial State	1	2	3
	4		5
	7	8	6

Goal State	1	2	3
	4	5	6
	7	8	

# Introducing the A\* Search Algorithm

A\* (pronounced "A-star") is a powerful and widely used pathfinding and graph traversal algorithm.

## Key Features:

- **Informed Search:** It uses a heuristic to guide its search, making it much more efficient than uninformed methods like Breadth-First Search (BFS).
- **Optimality:** If the heuristic is "admissible" (never overestimates the cost), A\* is guaranteed to find the shortest path.
- **Completeness:** It will always find a solution if one exists.

A\* is a perfect fit for the 8-puzzle because it intelligently explores the most promising moves first, avoiding a brute-force search of all possible states.

# The Core of A\*: The Evaluation Function

A\* decides which path to explore next based on the following evaluation function for a node (or state)  $n$ :

## Evaluation Function

$$f(n) = g(n) + h(n)$$

Where:

- $g(n)$ : The actual cost of the path from the start node to node  $n$ .
  - For the 8-puzzle, this is simply the number of moves made so far.
- $h(n)$ : The **heuristic** estimated cost from node  $n$  to the goal.
  - This is the "intelligent" part of A\*. It's an educated guess.

The algorithm always expands the node with the **lowest**  $f(n)$  value.

# Heuristics for the 8-Puzzle

A good heuristic is crucial for A\*'s performance. Two common admissible heuristics for the 8-puzzle are:

## ① Number of Misplaced Tiles

- Simple: Count the number of tiles that are not in their goal position.
- Fast to compute, but less accurate.

## ② Manhattan Distance (Most Common)

- For each tile, sum the number of horizontal and vertical moves required to get it to its goal position.
- More accurate than misplaced tiles, leading to a more efficient search.

### Example: Manhattan Distance

For tile '6' in the initial state, its goal is one step left and one step up. Its Manhattan distance is  $1 + 1 = 2$ .

# A\* Algorithm Step-by-Step

## Data Structures Needed:

- **Open List (Priority Queue):** Stores nodes that have been generated but not yet visited. Nodes are ordered by their  $f(n)$  value.
- **Closed List (Set or Hash Table):** Stores nodes that have already been visited to avoid cycles and redundant work.

## The Process:

- ① Create the start node and add it to the Open List.
- ② While the Open List is not empty:
  - ① Pop the node with the smallest  $f(n)$  value from the Open List. Let's call it *current*.
  - ② If *current* is the goal state, we're done! Reconstruct the path.
  - ③ Add *current* to the Closed List.
  - ④ Generate all valid successor nodes of *current*.
  - ⑤ For each successor:
    - If it's already in the Closed List, ignore it.
    - Calculate its  $g(n)$  and  $h(n)$  values.
    - If it's not in the Open List, add it.
- ③ If the Open List becomes empty, no solution exists.

# Your Goal

Your objective is to complete a Python script that finds the shortest solution for the 8-puzzle problem using the A\* search algorithm.

You have been provided with a script containing:

- A 'Node' class to represent puzzle states.
- Several helper functions.
- A complete testing framework.

## Your Task

You must implement the logic for three key functions that are currently empty or incomplete.

# Task 1: Implement the Heuristic

**Function to Complete:** `calculate_manhattan_distance`

## What it does

This function is the heuristic ( $h(n)$ ) for our A\* algorithm. It must calculate the total Manhattan distance for the given puzzle 'state'.

## Input:

- `state`: A 3x3 list of lists (the current board).
- `goal_positions`: A dictionary mapping each tile to its goal '(row, col)'.

## Output:

- An integer sum of the Manhattan distances for all tiles.

## Hint:

- Loop through each cell of the 'state'.
- For each tile, find its distance to its goal position:  
`abs(row - goal_row) + abs(col - goal_col)`.
- Sum these distances.



## Task 2: Generate Successor States

**Function to Complete:** `generate_successors`

### What it does

This function finds all valid next states by swapping the blank tile ('0') with an adjacent tile.

### Input:

- `state`: A 3x3 list of lists (the current board).

### Output:

- A list of new 3x3 states, one for each valid move.

### Hint:

- Find the location of the blank tile ('0').
- For each potential move (up, down, left, right), check if it's within the grid.
- If a move is valid, create a **deep copy** of the state, perform the swap, and add the new state to your list.

## Task 3: Implement the A\* Solver

**Function to Complete:** `solve`

### What it does

This is the main A\* algorithm logic. It uses a priority queue ('open\_list') and a visited set ('closed\_set') to find the optimal path.

### Process to Implement:

- ① Start a 'while' loop that runs as long as the 'open\_list' is not empty.
- ② Pop the node with the lowest f-score from the 'open\_list'.
- ③ **Goal Check:** If it's the goal state, reconstruct and return the path.
- ④ Add the current node to the 'closed\_set'.
- ⑤ Generate all successors for the current node.
- ⑥ For each successor:
  - If it's in the 'closed\_set', ignore it.
  - Otherwise, create a new 'child\_node' and add it to the 'open\_list'.

# How to Test Your Code

Once you have completed the functions, run the script from your terminal:

```
python your_script_name.py
```

The testing framework will automatically evaluate your solution against several test cases.

## You will see:

- **PASSED:** Correct!
- **FAILED:** Incorrect.

## A passing grade requires:

- Finding the optimal (shortest) path.
- Correctly identifying unsolvable puzzles.

# Good Luck!